

BT5110 Data Management and Warehousing

Tutorial 2: Simple and Algebraic Queries

Mark Meng Huasong

School of Computing
National University of Singapore

30 Aug - 3 Sep, 2021

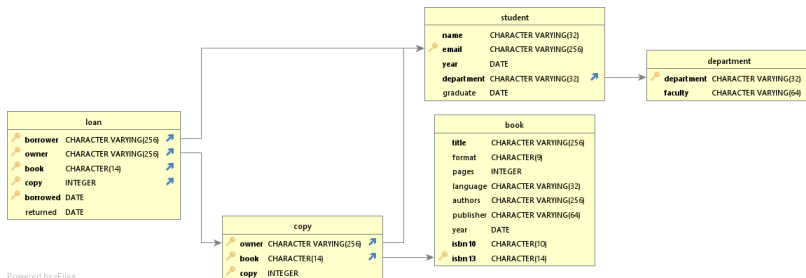


All the materials within presentation slides are protected by copyrights.
It is forbidden by NUS to upload these materials to the Internet.

Quick Recap: End of Last Tutorial

What we have done in the last week:

- (1) Create 4 tables and populate data;
- (2) Update all 'CS' with 'Computer Science' in department column;
- (3) Drop available from table copy;
- (4) Create a separate table department and migrate faculty from table student to table department.



(plotted by DbVisualizer)

Question 1 (a-c)

- (a) Print the different departments.
- (b) Print the different departments in which students are enrolled.
- (c) Let us check the integrity of the data. Print the emails of the students who borrowed or lent a copy of a book before they joined the university. There should not be any. Use a simple query.

Solution 1 (a, b)

(a) Print the different departments.

```
1  SELECT d.department FROM department d;
```

13 row are returned. Notice that the query does not require **DISTINCT** to eliminate duplicates. Duplicates are guaranteed not to occur because department is the **PRIMARY KEY** of the table department.

(b) Print the different departments in which students are enrolled.

```
1  SELECT DISTINCT s.department FROM student s;
```

13 row are returned. There could be departments in which no student is enrolled. This is the case of the department of Undecidable Computations. We need to look into the student table.

Do these two questions return the exactly SAME outputs?

Solution 1 (a, b) Cont.

Notice

The outputs of these two queries have the same contents but with different orders.

```
BT5110=# select d.department from department d;  
department
```

```
-----  
CS  
History  
Physics  
Geography  
Language  
Economics  
Biology  
EE  
Math  
IS  
CE  
ME  
Chemistry  
(13 rows)
```

```
BT5110=# select DISTINCT s.department from student s;  
department
```

```
-----  
EE  
CS  
ME  
Language  
Math  
CE  
IS  
Physics  
Economics  
Geography  
Biology  
Chemistry  
History  
(13 rows)
```

Extra: We can add “ORDER BY d.department ASC” in 1(a) and “ORDER BY s.department ASC” in 1(b) to make these two outputs same in ordering.

Solution 1 (c)

(c) Let us check the integrity of the data. Print the emails of the students who borrowed or lent a copy of a book before they joined the university. **There should not be any.** Use a simple query.

Solution: Don't forget to consider both *borrowing* and *lending* scenarios.

```
1  SELECT DISTINCT s.email FROM loan l, student s
2  WHERE (s.email = l.borrower AND l.borrowed < s.year)
3  OR (s.email = l.owner AND l.borrowed < s.year);
```

Question 1 (d)

For each copy that has been borrowed and returned, print the **duration** of the loan. Order the results in **ascending** order of the ISBN13 and **descending** order of duration.

How can the duration be derived? Can we use “returned - borrowed AS duration”?

```
1  SELECT book, returned - borrowed AS duration
2  FROM loan
3  ORDER BY book ASC, duration DESC;
```

Any issue in the code above?

Solution 1 (d)

Let's manually exclude NULL values in returned column:

```
1  SELECT book, returned - borrowed + 1 AS duration
2  FROM loan
3  WHERE NOT (returned ISNULL)
4  ORDER BY book ASC, duration DESC;
```

Notice

ASC is the default, but it is strongly recommended to indicate it for clarity.

Result: 4871 rows are returned (\neq number of rows in table loan).

Is this result good enough? We do have loan records with NULL returned value!

Solution 1 (d) Cont.

Notice that the duration can be **null** if the book has not been returned yet. For a complete answer, you need to calculate the duration until a specific date (e.g., December 31, 2010) to include the books that have not been returned yet.

```
1  SELECT book, ((CASE WHEN returned ISNULL
2      THEN '2010-12-31'
3      ELSE returned END) - borrowed + 1) AS duration
4  FROM loan
5  ORDER BY book ASC, duration ASC;
```

Result: 4976 rows are returned (= number of rows in table loan)

You can also use a Postgres reserved command “current_date” to obtain the date of today (by doing this, some duration values will be huge).

Question 1 (e)

For each loan of a book published by Wiley that has not been returned, print the title of the book, the name and faculty of the owner and the name and faculty of the borrower. Use **CROSS JOIN**.

Let's do 4 steps to form your query as follows:

Step (1) FIND all output columns;

Step (2) IDENTIFY all value constraints;

Step (3) CONFIRM all tables involved;

Step (4) CONNECT tables involved (w. necessary relation constraints).

*** LIVE DEMO ***

Solution 1 (e)

We join primary keys and foreign keys to *stitch* tables together properly.

```
1  SELECT b.title,  
2     s1.name AS ownername,  
3     d1.faculty AS ownerFaculty,  
4     s2.name AS borrowername,  
5     d2.faculty AS borrowerfaculty  
6  FROM loan l, book b, copy c,  
7     student s1, student s2,  
8     department d1, department d2  
9  WHERE c.book = b.ISBN13      ①  
10 AND c.book = l.book         ②  
11 AND c.copy = l.copy         ③  
12 AND c.owner = l.owner       ④  
13 AND l.owner = s1.email      ⑤  
14 AND l.borrower = s2.email   ⑥  
15 AND s1.department = d1.department ⑦  
16 AND s2.department = d2.department ⑧  
17 AND b.publisher = 'Wiley'   A  
18 AND l.returned ISNULL;      B
```

10 rows are returned.

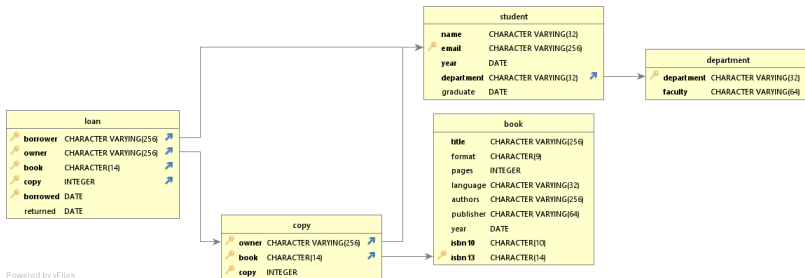
Solution 1 (e) Cont.

You can omit the table copy and the copy column since the existence of the corresponding rows and values is guaranteed by design and by the foreign and primary key constraints.

```
1  SELECT b.title,
2     s1.name AS ownername,
3     d1.faculty AS ownerFaculty,
4     s2.name AS borrowername,
5     d2.faculty AS borrowerfaculty
6  FROM loan l, book b,
7     student s1, student s2,
8     department d1, department d2
9  WHERE l.book = b.ISBN13
10     AND l.owner = s1.email
11     AND l.borrower = s2.email
12     AND s1.department = d1.department
13     AND s2.department = d2.department
14     AND b.publisher = 'Wiley'
15     AND l.returned ISNULL;
```

Question 2 (a)

For each loan of a book published by Wiley that has not been returned, print the title of the book, the name and faculty of the owner and the name and faculty of the borrower. Use **INNER JOIN**.



Powered by yFiles

This is the same question with Q1(e), but requires you to use **INNER JOIN**.

Solution 2 (a)

Let's convert the solution of Q1(e) now!

```
1  SELECT b.title,
2     s1.name AS ownername,
3     d1.faculty AS ownerFaculty,
4     s2.name AS borrowername,
5     d2.faculty AS borrowerfaculty
6  FROM loan l, book b,
7     student s1, student s2,
8     department d1, department d2
9  WHERE l.book = b.ISBN13
10     AND l.owner = s1.email
11     AND l.borrower = s2.email
12     AND s1.department = d1.department
13     AND s2.department = d2.department
14     AND b.publisher = 'Wiley'
15     AND l.returned ISNULL;
```

```
1  SELECT b.title,
2     s1.name AS ownername,
3     d1.faculty AS ownerFaculty,
4     s2.name AS borrowername,
5     d2.faculty AS
        borrowerfaculty
6  FROM loan l
7     INNER JOIN _____ ON _____
8     INNER JOIN _____ ON _____
9     INNER JOIN _____ ON _____
10    INNER JOIN _____ ON _____
11    INNER JOIN _____ ON _____
12  WHERE b.publisher = 'Wiley'
13     AND l.returned ISNULL;
```

Solution 2 (a) Cont.

You can omit the table copy and the copy column since the existence of the corresponding rows and values is guaranteed by design and by the foreign and primary key constraints.

```
1  SELECT b.title,
2     s1.name AS ownername,
3     d1.faculty AS ownerFaculty,
4     s2.name AS borrowername,
5     d2.faculty AS borrowerfaculty
6  FROM loan l
7     INNER JOIN book b ON l.book=b.ISBN13
8     INNER JOIN student s1 ON l.owner = s1.email
9     INNER JOIN student s2 ON l.borrower = s2.email
10    INNER JOIN department d1 ON s1.department = d1.department
11    INNER JOIN department d2 ON s2.department = d2.department
12 WHERE b.publisher = 'Wiley'
13    AND l.returned ISNULL;
```

10 rows are returned.

Solution 2 (a) Cont.

The code below is without omitting table copy. The simplified version is shown on the previous slide.

```
1  SELECT b.title ,
2     s1. name AS ownername ,
3     d1. faculty AS ownerFaculty ,
4     s2. name AS borrowername ,
5     d2. faculty AS borrowerfaculty
6  FROM loan l
7  INNER JOIN book b ON l. book = b. ISBN13
8  INNER JOIN copy c ON c. book = l. book
9  AND c. copy = l. copy AND c. owner = l. owner
10 INNER JOIN student s1 ON l. owner = s1. email
11 INNER JOIN student s2 ON l. borrower = s2. email
12 INNER JOIN department d1 ON s1. department = d1. department
13 INNER JOIN department d2 ON s2. department = d2. department
14 WHERE b. publisher = 'Wiley'
15 AND l. returned ISNULL ;
```

Question 2 (b)

Print the emails of the different students who borrowed or lent a copy of a book on the day that they joined the university. Use an algebraic query.

Set operations (UNION, INTERSECT and EXCEPT) could help to address this question.

Solution 2 (b)

```
1  SELECT  s.email
2  FROM    loan l, student s
3  WHERE   s.email = l.borrower AND l.borrowed = s.year
4  UNION
5  SELECT  s.email
6  FROM    loan l, student s
7  WHERE   s.email = l.owner AND l.borrowed = s.year;
```

19 rows are returned.

DISTINCT is not needed because UNION eliminates duplicates (so do INTERSECT, EXCEPT and MINUS).

Question 2 (b) Cont.

There is an alternative (simple) way to write the query, without SET operations.

The corresponding simple query is generally preferable.

```
1  SELECT DISTINCT s.email
2  FROM loan l,  student s
3  WHERE (s.email = l.borrower OR s.email = l.owner)
4      AND l.borrowed = s.year;
```

Notice

The simple query requires an explicit DISTINCT.

Question 2 (c)

Print the emails of the different students who borrowed and lent a copy of a book on the day that they joined the university. Use an algebraic query.

Which set operation (UNION, INTERSECT and EXCEPT) should we use for this question?

Solution 2 (c)

```
1  SELECT  s.email
2  FROM loan l, student s
3  WHERE s.email = l.borrower AND l.borrowed = s.year
4  INTERSECT
5  SELECT  s.email
6  FROM loan l, student s
7  WHERE s.email = l.owner AND l.borrowed = s.year;
```

4 rows are returned.

Note that the corresponding simple query is more *complicated*. It needs *two* loan tables.

```
1  SELECT DISTINCT s.email
2  FROM loan l1, loan l2, student s
3  WHERE s.email = l1.borrower AND l1.borrowed = s.year
4  AND s.email = l2.owner AND l2.borrowed = s.year;
```

Question 2 (d)

Print the emails of the students who borrowed but did not lend a copy of a book on the day that they joined the university. Use an algebraic query.

Which set operation should we use for this question?

Solution 2 (d)

```
1  SELECT  s.email
2  FROM loan l, student s
3  WHERE s.email = l.borrower AND l.borrowed = s.year
4  EXCEPT
5  SELECT  s.email
6  FROM loan l, student s
7  WHERE s.email = l.owner AND l.borrowed = s.year;
```

9 rows are returned.

There is no corresponding simple query. We would need to use nested or aggregate queries for this type of questions.

Question 2 (e)

Print the ISBN13 of the books that have never been borrowed. Use an algebraic query.

Solution 2 (e)

Solution:

```
1  SELECT b.ISBN13
2  FROM book b
3  EXCEPT
4  SELECT l.book
5  FROM loan l;
```

0 rows are returned.

or, using an *OUTER JOIN*, which introduces NULL values,

```
1  SELECT b.ISBN13
2  FROM book b LEFT OUTER JOIN loan l ON b.isbn13 = l.book
3  WHERE l.book ISNULL;
```

There is no corresponding simple query. We would need to use nested or aggregate queries for this type of questions.

For any further question, please feel free to email me:

huasong.meng@u.nus.edu

Copyright 2021 Mark H. Meng. All rights reserved.